



# CS207



**SYSTEMS DEVELOPMENT FOR  
COMPUTATIONAL SCIENCE**

**AKA**

**COMPUTER SCIENCE FOR SCIENTISTS**

**CONCEIVED WITH LOVE @ IACS**

## THIS TALK:

1. WHY DO THIS COURSE
2. WHAT'S IN THE COURSE
3. LOOK AT SCHEDULE AND SOME PRACTICALITIES


**WHY DO IT?**

# SPARK<sup>1</sup>


- > 'BIG DATA' ANALYSIS
  - > IMMUTABLE DATA TABLES, CALLED RDDEES
  - > RDDS PARTITIONED BETWEEN PROCESSES
- > LAZY EVALUATION OF TRANSFORMATIONS BETWEEN RDDEES
  - > BCOZ GRAPH BASED REPRESENTATION OF RDD TRANSFORMATIONS

<sup>1</sup> DIAGS FROM [HTTPS://WWW.CS.CMU.EDU/~PAVLO/COURSES/FALL2013/STATIC/SLIDES/SPARK.PDF](https://www.cs.cmu.edu/~pavlo/courses/fall2013/static/slides/spark.pdf)


## Narrow Dependencies:



map, filter




union




join with inputs  
co-partitioned

## Wide Dependencies:




groupByKey



join with inputs not  
co-partitioned

# Job Scheduling







- FAULT RECOVERY USING LINEAGE FROM GRAPH
- FOR NARROWS, ONLY PARTIAL FAULT RECOVERY NEEDED (SEE GRAPH)
- KEY IDEAS: IMMUTABILITY GUARANTEES PARTICULAR DATASET AT ANY STAGE
- KEY IDEAS: COARSE GRAINED COMPUTATION REPRESENTED AS GRAPH ALLOWS SCHEDULING AND RECOVERY
- GRAPH DEPENDENCY SORTING: MAKEFILES, THEANO, SPARK, DASK, ETC, ETC



# ANOTHER EXAMPLE: ASYNCHRONOUS SGD



(FROM ANDREW NG'S COURSERA COURSE)

TO MINIMIZE  $Q(w) = \sum_{i=1}^n Q_i(w)$  (FROM WIKIPEDIA)


- CHOOSE AN INITIAL VECTOR OF PARAMETERS AND LEARNING RATE.

- REPEAT UNTIL AN APPROXIMATE MINIMUM IS OBTAINED:

1. RANDOMLY SHUFFLE EXAMPLES IN THE TRAINING SET.

2. FOR  $i=1..N$ , DO:

- $w_{new} = w - \eta \nabla Q_i(w)$

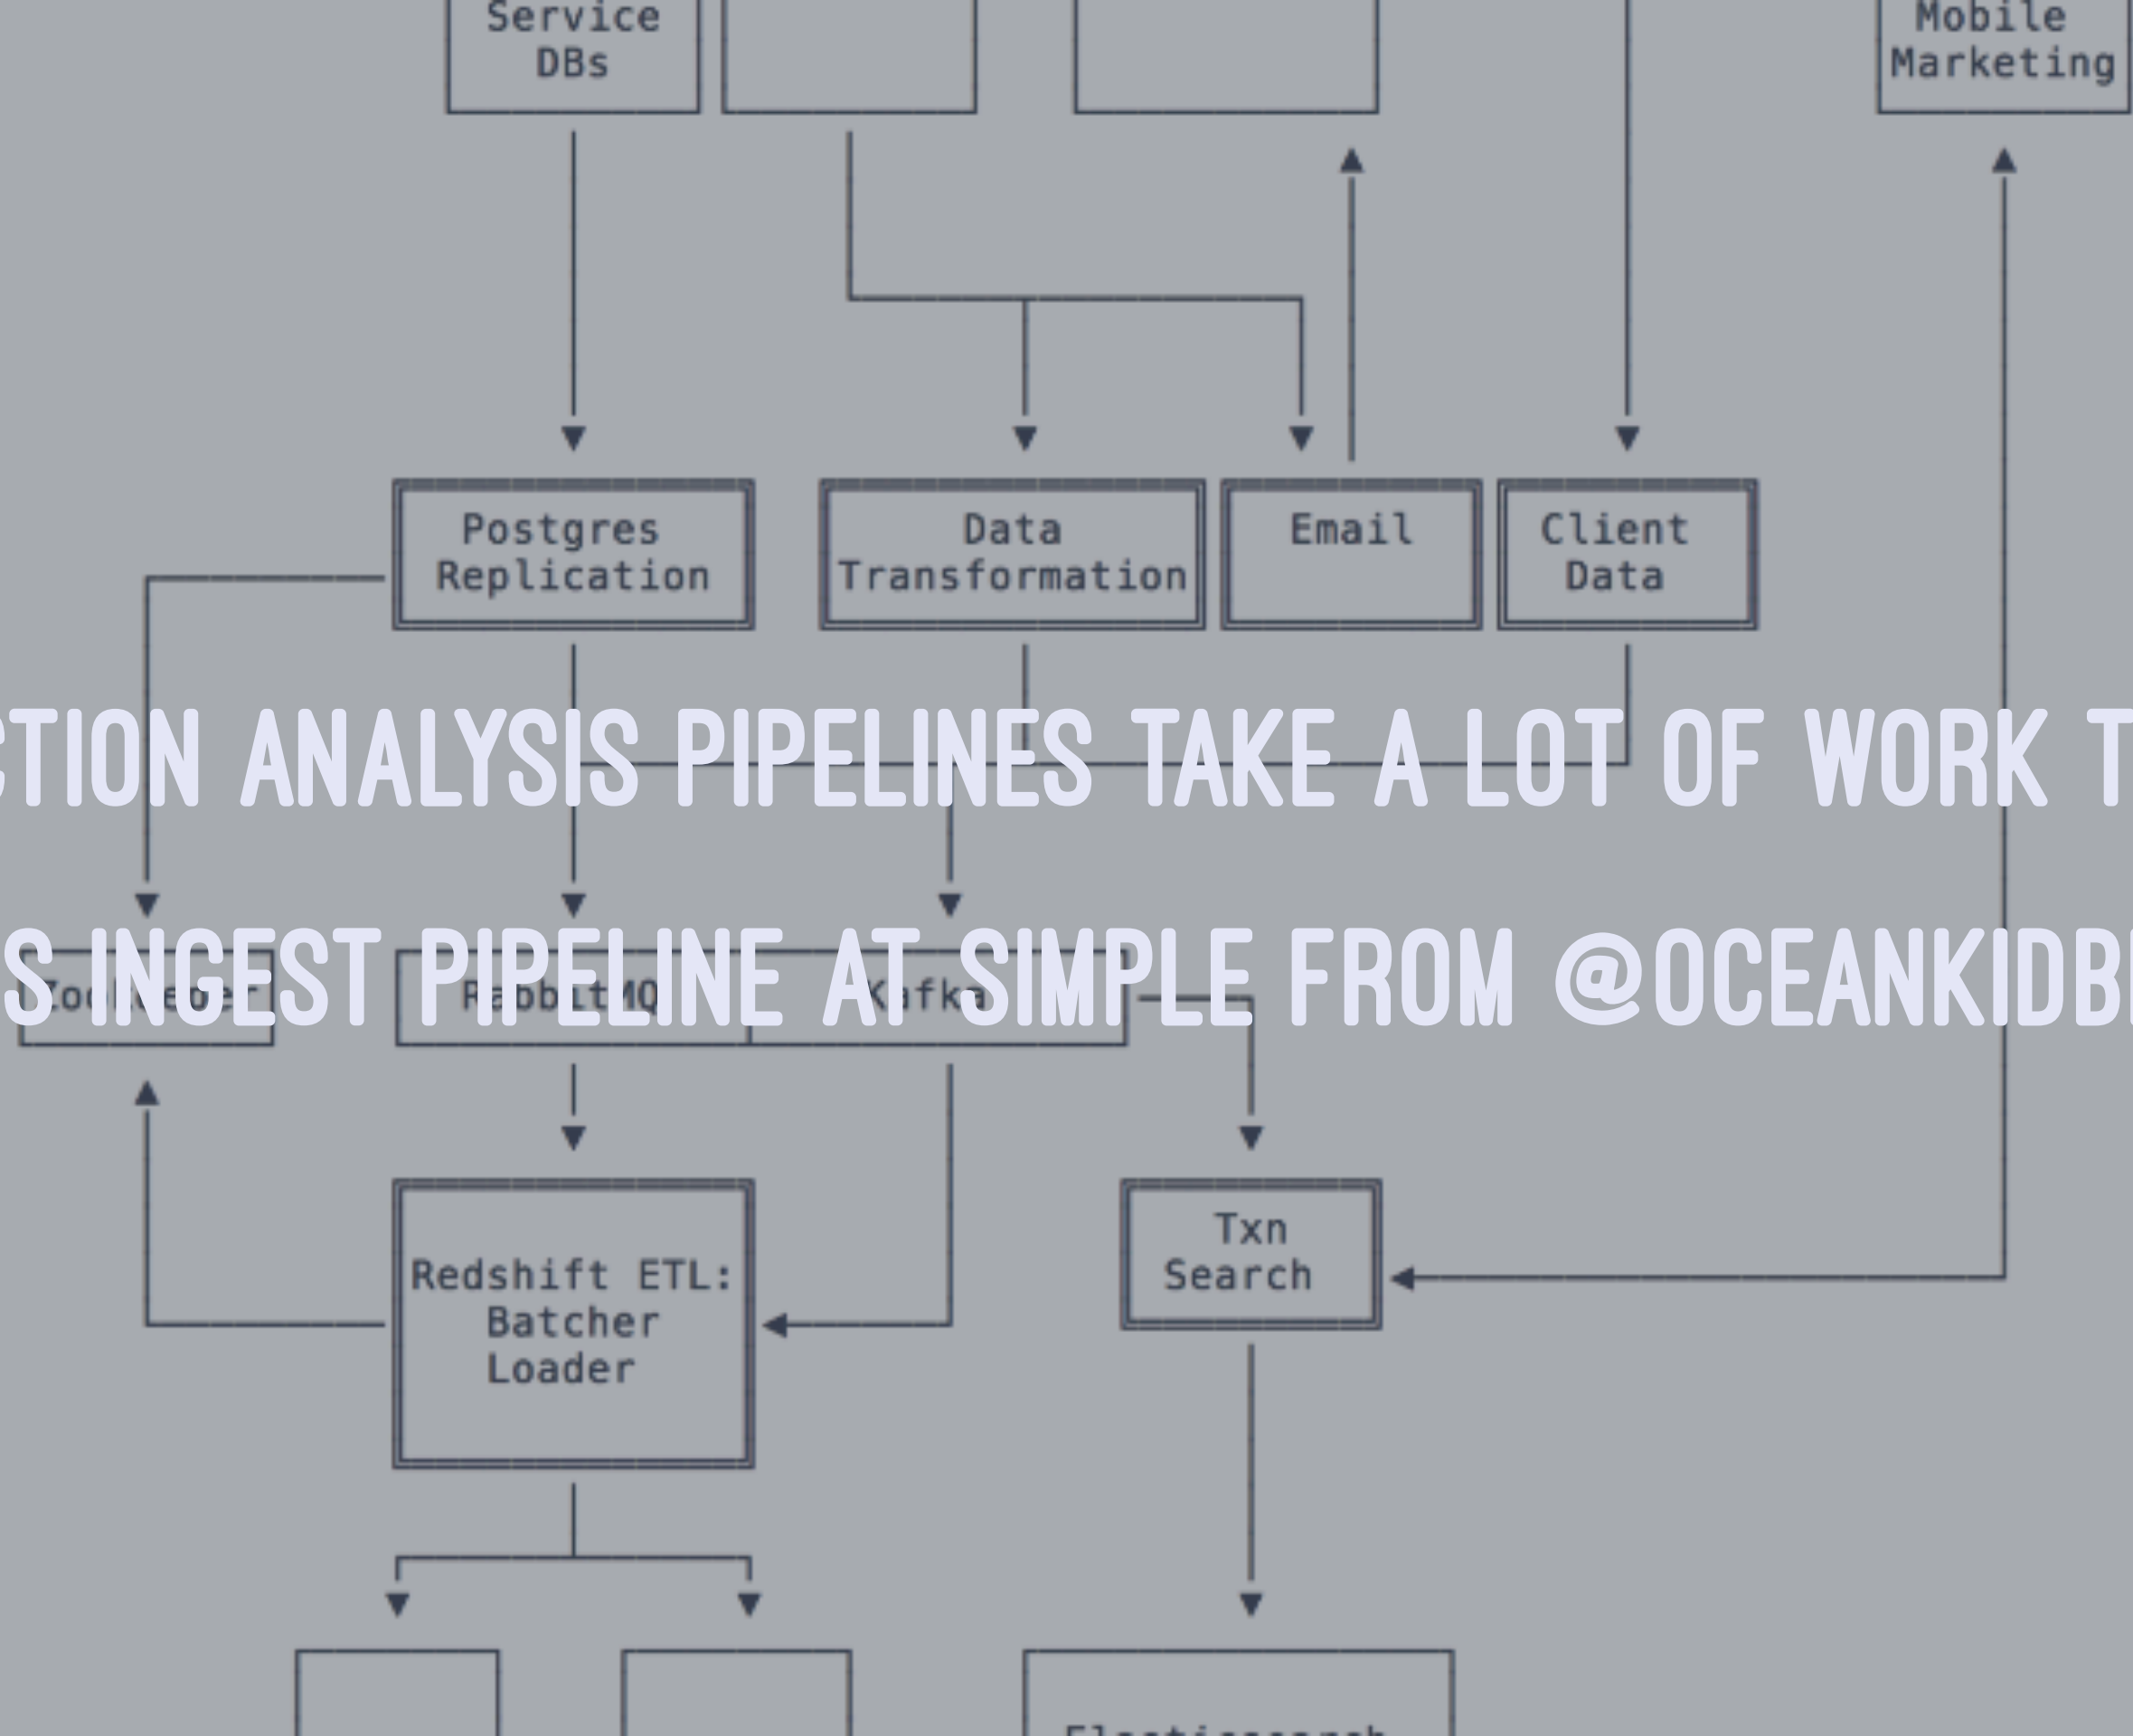


- > SEE DEAN ET.AL 2012
- > RUN COPY OF MODEL ON EACH BATCH (MULTIPLE BATCHES)
- > COMMUNICATE PARAMETERS WITH SHARDED PARAMETER SERVER ASYNCHRONOUSLY (MULTIPLE PARAMETER SHARDS)

- ROBUST TO MACHINE FAILURE
    - BUT MORE STOCHASTIC
- NOT GUARANTEED UPDATED PARAMETERS
- IN PRACTICE RELAXING CONSISTENCY WORKS!
  - BETTER FOR MORE LOCAL NEURAL NETS

**SCALING FROM RESEARCH TO PRODUCTION CAN BE COMPLEX:**  
**READ [HTTPS://OPENAI.COM/BLOG/INFRASTRUCTURE-FOR-DEEP-LEARNING/](https://openai.com/blog/infrastructure-for-deep-learning/)**  
**NEED TO PROVISION CLUSTERS OF GPUS/CPUS STARTING FROM 1 MACHINE TO MANY...**

PRODUCTION ANALYSIS PIPELINES TAKE A LOT OF WORK TOO!  
ANALYSIS INGEST PIPELINE AT SIMPLE FROM @OCEANKIDBILLY



**YOU MAY, ADDITIONALLY, HAVE TO CREATE:**

- > INTERNAL DASHBOARDS AND APPS FOR BUSINESS DECISIONS**
- > THUS WILL HAVE TO BE FAMILIAR WITH WEB TECHNOLOGIES**
  - > KNOW HOW TO USE VIRTUALIZED AND CONTAINERIZED INFRASTRUCTURE**
- > PRODUCE AND CONSUME INTERNAL AND EXTERNAL API ENDPOINTS**



**HOW TO START?**



A COURSE WITH 3 SIMULTANEOUS THEMES...

1. ENGINEERING A TIGHT  
SHIP

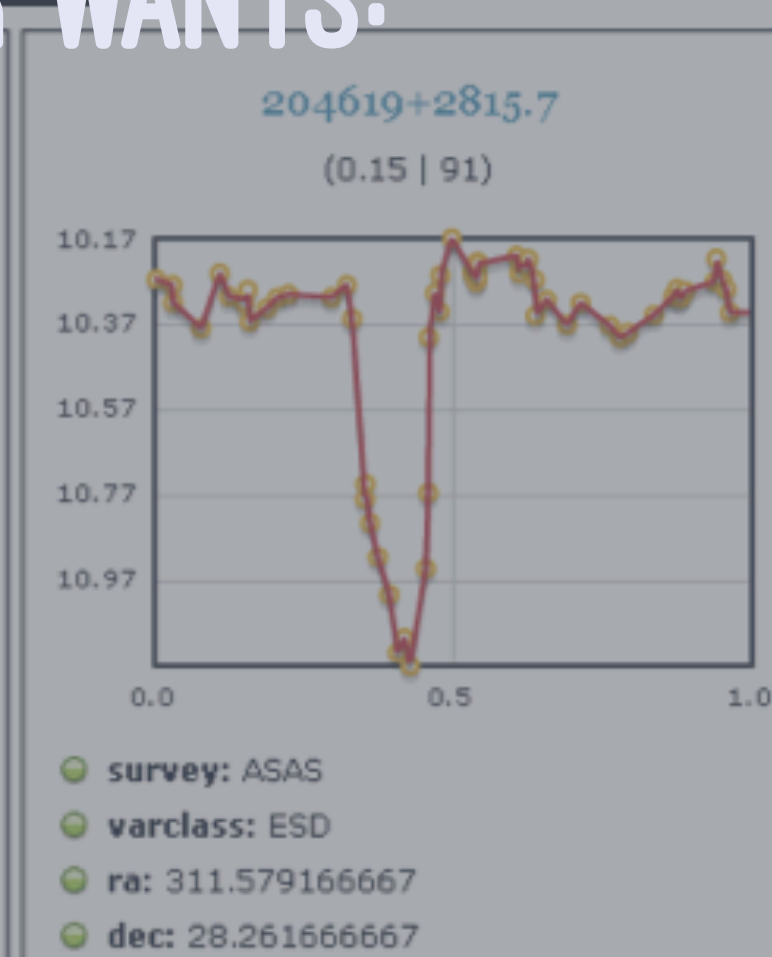
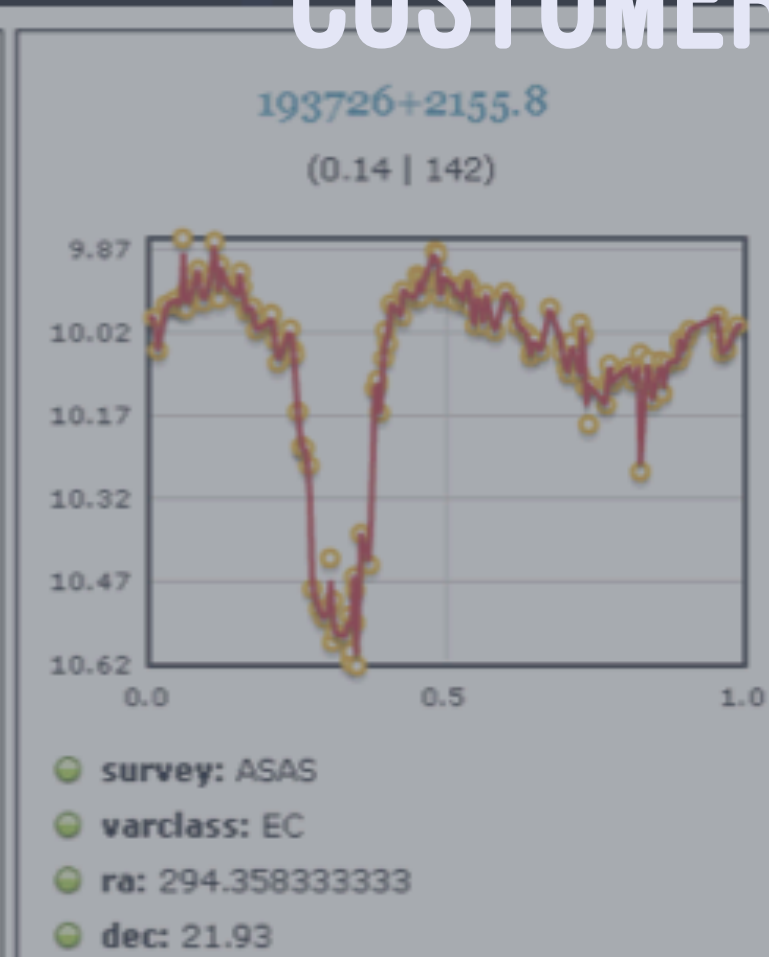
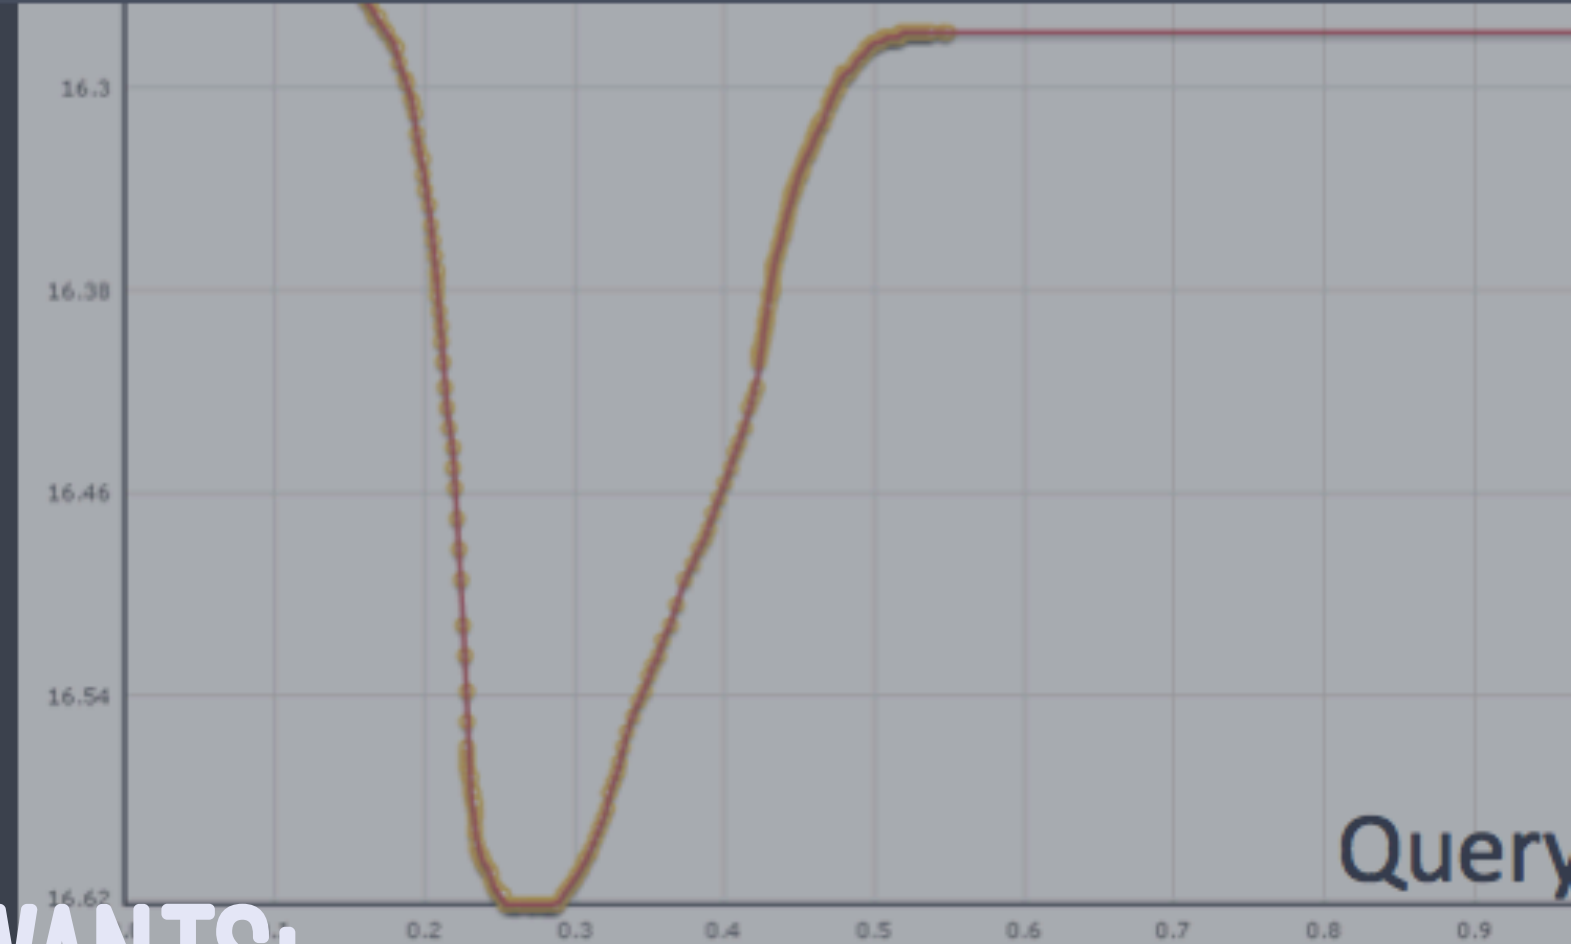
2. WHAT'S IN THE  
LANGUAGE

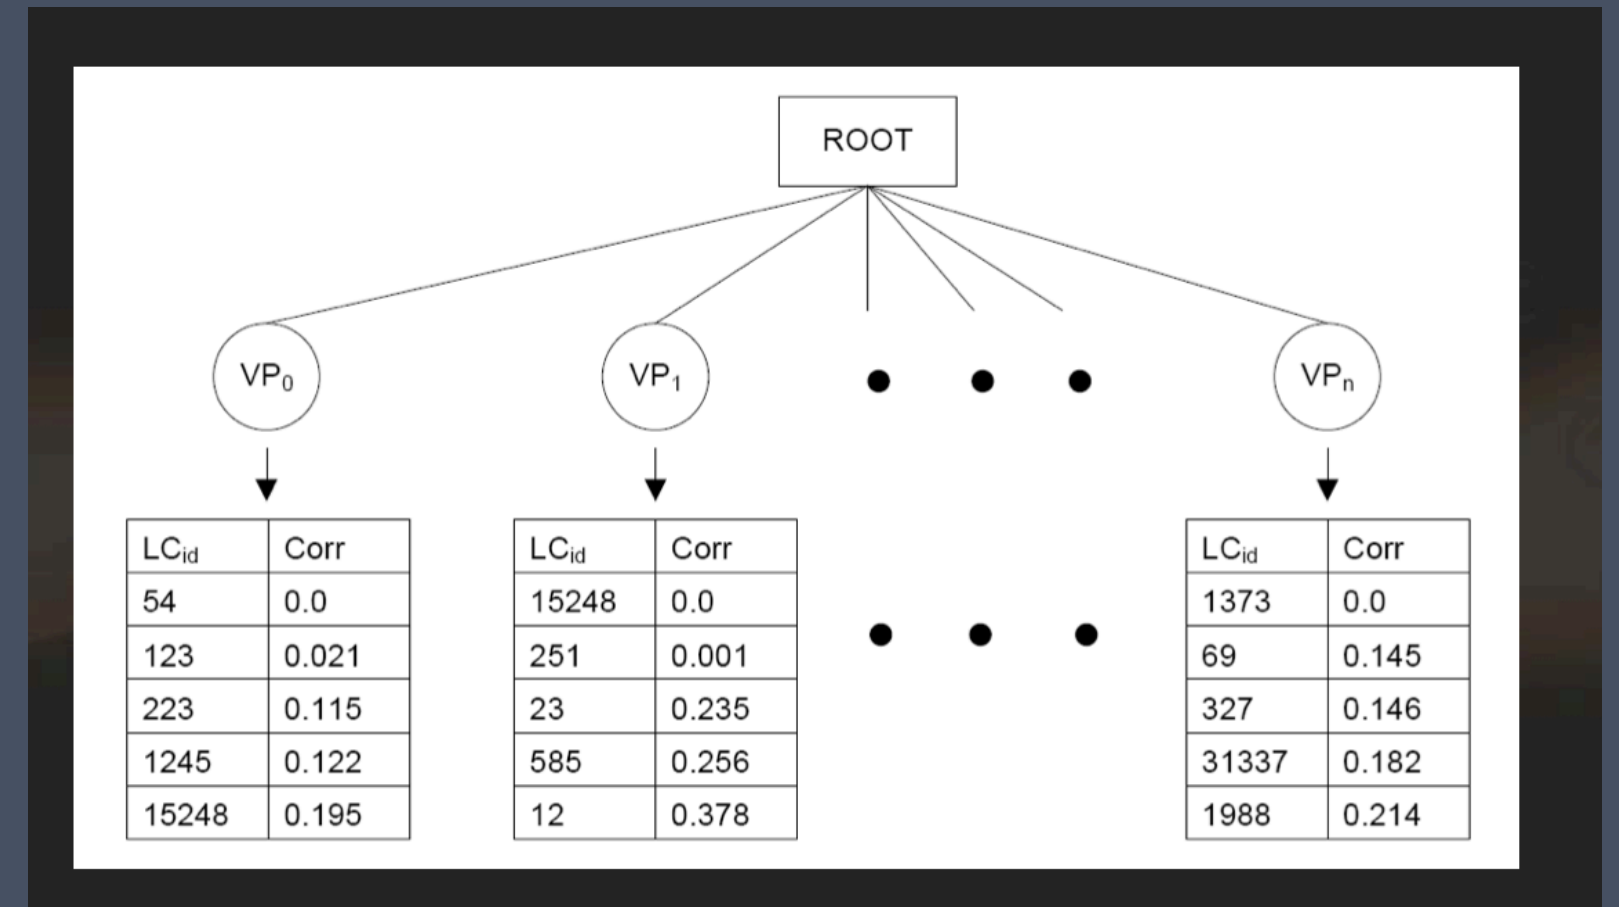

3. ALGORITHMS FIND  
DATA.

# Problem.

- Searching Light Curves
- 100-1000 Data-points

CUSTOMER WANTS:





KNN SEARCH IN A LARGE SPACE, LEADS TO AN INDEX...

# I. A TIME SERIES LIBRARY

(WITH OPERATIONS BETWEEN, AND ALGORITHMS ON TIME SERIES)

- > PYTHON

- > EXECUTION MODEL, PROGRAMMING PARADIGMS, OBJECT MODEL

- > SEQUENCES, ITERATORS, LISTS, ARRAYS, TREES, HASHES

- > TESTS, DEBUGGING, CI, AUTOMATING

# II. A TIME SERIES DATABASE

(WITH INDEXING, SIMILARITY QUERIES, QUERY LANGUAGE)

- WRITING A CLI, A REPL, A DSL, WEB ACCESS
- APPROPRIATE ON-DISK DATA STRUCTURES FOR THE INDEXES AND TS
  - USAGE OF A DATABASE



# III. A REST INTERFACE AND UI TO THESE AND MORE

(THE **MORE** IS YOUR CHOICE, BUT YOU MUST DELIVER A DEMO)

- > YOU WRITE AN EXTENSION OF **YOUR CHOICE**
- > YOU PACKAGE EVERYTHING UP IN ONE OR MORE SERVERS,  
ACCESSSED BY A  
WEB API AND INTERFACE, AND PROVIDE A DEMO



# HOW?

- ALL WORK IS ON GITHUB.
- ONE HOMEWORK A WEEK, TO BE SUBMITTED INDIVIDUALLY
- YOU WILL WORK IN GROUPS OF 3-4. YOU CAN COLLABORATE AS MUCH AS YOU WANT WITHIN THE GROUP (INCLUDING HOMEWORK, BUT YOU MUST WRITE UP YOUR OWN HOMEWORK).
  - ONE REPO PER GROUP FOR YOUR PROJECT.
  - NO EXAMS

# HOW (CONTD)

- > THE BASIC PART OF THE PROJECT (I AND II) WILL BE DONE BY ALL GROUPS.
- > ONCE A GROUP WILL IMPLEMENT A FEATURE FOR ANOTHER GROUP
- > GROUPS WILL DO DIFFERENT THINGS FOR THE ELECTIVE PART AND REST API, WEB UI, DATABASE

# WHEN?

- > MW 2.30PM – 4PM LECTURE PIERCE 209
- > FRIDAYS 3PM–4.30PM (LAB) PIERCE 301

# GRADE

- > HOMEWORK (40%)
- > PROJECT (45%)
- > PROJECT X-DEVEL (7%)
- > PARTICIPATION (8%): IN DISCUSSIONS, COMMITS, PEER REVIEW AND CODE REVIEW EXERCISE IN M1 OR M2

YOU WILL PEER-REVIEW EACH OTHER.

WE'LL DO PRACTICALITIES IN LAB 1 THIS FRIDAY. BUT IF YOU WANT TO GET STARTED:

- > INSTALL ANACONDA PYTHON (PYTHON 3.5)
- > CREATE A GITHUB ACCOUNT. WITH A REPO CALLED `cs207work`
- > INSTALL A GIT-CLIENT. ON WINDOWS INSTALL GIT-BASH.
- > YOU MAY OPTIONALLY WANT TO INSTALL THE GITHUB CLIENT.
  - > CHOOSE A CODE EDITOR: I LIKE ATOM.

- 
- A white unicorn with a single golden horn is running across a green grassy field. In the background, there is a wooden fence and a line of trees. The image has a dark blue overlay on the left and right sides.
- > FUN IS GUARANTEED!
  - > LOTS OF COOL STUFF TO LEARN
  - > LEARN TO LEARN, NOT TO GET A GRADE

# UNICORN LEVEL 1 UNLOCKED.